

Robotic Navigation and Exploration



HW4: Deep Reinforcement Learning with PyTorch

Min-Chun Hu anitahu@cs.nthu.edu.tw
CS, NTHU

Navigation Environment

Navigation Environment

- Path following




Navigation Environment

- **State:** future positions + past positions + past yaws ($s \in \mathbb{R}^{14}$)
- **Action:** current angular velocity ($a \in \mathbb{R}$)
- **Reward:** $r = w_d r_d + w_y r_y + r_p$

Distance reward: $r_d = \exp(-0.1 \|p - p'\|^2)$

Yaw reward: $r_y = \exp(-0.1 \|\theta - \theta'\|^2)$

Progress reward: $r_p = \begin{cases} 0.1 & , \text{ if progress is positive} \\ 0 & , \text{ if progress is 0} \\ -1 & , \text{ if progress is negative} \end{cases}$



- **Termination:**

100% progress or reach 400 steps

Note:

The value range of action is normalized to $[-1, 1]$. If the input action value is out of range, the value will be clipped.

PPO Algorithm

PPO Algorithm

Algorithm 1 PPO, Actor-Critic Style

```
for iteration=1,2,... do  
  for actor=1,2,...,  $N$  do  
    Run policy  $\pi_{\theta_{\text{old}}}$  in environment for  $T$  timesteps  
    Compute advantage estimates  $\hat{A}_1, \dots, \hat{A}_T$   
  end for  
  Optimize surrogate  $L$  wrt  $\theta$ , with  $K$  epochs and minibatch size  $M \leq NT$   
   $\theta_{\text{old}} \leftarrow \theta$   
end for
```


Run an episode to
collect data

Train the model using
the collected data

PPO Algorithm


- Actor Network:

$$\pi_{\theta}(a|s)$$



- Critic Network:

$$V_{\omega}(s)$$



PPO Algorithm

- Value Loss:

$$L(\omega) = \frac{1}{2} \sum_n [G_t - V_\omega(s_t^{(n)})]^2$$

- Policy Gradient Loss:


$$r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$$

$$L^{CLIP}(\theta) = E_{(s_t, a_t) \sim \pi_{\theta_{old}}} \left[\min \{ \underline{r_t(\theta) A^{\theta_{old}}(s_t, a_t)}, \quad \underline{\text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) A^{\theta_{old}}(s_t, a_t)} \} \right]$$

Policy Network & Value Network Construction

Diagonal Gaussian Distribution Module (model.py)

1. FixedNormal



$$N(x|\mu, \Sigma) = \frac{1}{(2\pi)^{(D/2)}} \frac{1}{|\Sigma|^{1/2}} \exp\left\{\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)\right\}$$

$$\ln p(X|\mu, \Sigma) = -\frac{ND}{2} \ln(2\pi) - \frac{N}{2} \ln |\Sigma| - \frac{1}{2} \sum_{n=1}^N (x_n - \mu)^T \Sigma^{-1} (x_n - \mu)$$

```
11 #Normal distribution module with fixed mean and std.
12 class FixedNormal(torch.distributions.Normal):
13     # Log-probability
14     def log_probs(self, actions):
15         return super().log_prob(actions).sum(-1)
16
17     # Entropy
18     def entropy(self):
19         return super().entropy().sum(-1)
20
21     # Mode
22     def mode(self):
23         return self.mean
24
```


Diagonal Gaussian Distribution Module (model.py)

2. DiagGaussian


$$\text{mean} = \begin{bmatrix} m_1 \\ m_2 \\ \vdots \end{bmatrix}, \quad \text{std} = \begin{bmatrix} \sigma_1 \\ \sigma_2 \\ \vdots \end{bmatrix} \Rightarrow \Sigma = \begin{bmatrix} \sigma_1^2 & 0 & \dots & 0 \\ 0 & \sigma_2^2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \sigma_k^2 \end{bmatrix}$$

```
25 #Diagonal Gaussian distribution
26 class DiagGaussian(nn.Module):
27     # Constructor
28     def __init__(self, inp_dim, out_dim, std=0.5):
29         super(DiagGaussian, self).__init__()
30
31         init_ = lambda m: init(
32             m,
33             nn.init.orthogonal_,
34             lambda x: nn.init.constant_(x, 0)
35         )
36         self.fc_mean = init_(nn.Linear(inp_dim, out_dim))
37         self.std = torch.full((out_dim,), std)
38
39     # Forward
40     def forward(self, x):
41         mean = self.fc_mean(x)
42         return FixedNormal(mean, self.std.to(x.device))
43
```


sigma = [1 0; 0 1]



sigma = [1 0; 0 4]



sigma = [4 0; 0 1]



Policy Network Module

- **PolicyNet** class has the following functions:

#Constructor

- `__init__(self, s_dim, a_dim, std)`

#Forward pass of nn.Module


- `forward(self, state, deterministic)`

#Forward pass for outputting action only

- `action_step(self, state, deterministic)`

#Evaluate log-prob. & entropy

- `evaluate(self, state, action)`




Policy Network Module (model.py)

- constructor (TODO 1)

(Hint: Use **nn.Sequential** & **DiagGaussian**)

```
44 #Policy network
45 class PolicyNet(nn.Module):
46     # Constructor
47     def __init__(self, s_dim, a_dim, std=0.5):
48         super(PolicyNet, self).__init__()
49
50         init_ = lambda m: init(
51             m,
52             nn.init.orthogonal_,
53             lambda x: nn.init.constant_(x, 0),
54             nn.init.calculate_gain('relu')
55         )
56         #TODO 1: policy network architecture
57         '''
58         self.main = ...
59         self.dist = ...
60         '''
```



Value Network Module


- **ValueNet** class has the following functions:

#Constructor

- `__init__(self, s_dim)`

#Forward pass of nn.Module

- `forward(self, state)`




Value Network Module (model.py)

- constructor (TODO 2)

(Hint: Use **nn.Sequential**)

```
93 #Value network
94 class ValueNet(nn.Module):
95     # Constructor
96     def __init__(self, s_dim):
97         super(ValueNet, self).__init__()
98
99         init_ = Lambda m: init(
100             m,
101             nn.init.orthogonal_,
102             Lambda x: nn.init.constant_(x, 0),
103             nn.init.calculate_gain('relu')
104         )
105         #TODO 2: value network architecture
106         '''
107         self.main = ...
108         '''
```



Environment Runner Construction

EnvRunner Class

- **EnvRunner** class has the following functions:

#Constructor

- `__init__(self, env, s_dim, a_dim, n_step, gamma, lamb, device)`

#Run n steps to get a batch

- `run(self, policy_net, value_net)`

#Record return & length

- `record(self)`

#Get current performance


- `get_performance(self)`

EnvRunner Class (env_runner.py)

- run (TODO 3)

(Hint: Use **policy_net** & **value_net**)


```
77 # Run n steps to get a batch
78 def run(self, policy_net, value_net):
79     #1. Run n steps
80     #-----
81     for step in range(self.n_step):
82         #self.states : (n_env, s_dim)
83         #actions      : (n_env, a_dim)
84         #self.dones   : (n_env)
85         #a_logps      : (n_env)
86         #values       : (n_env)
87         #rewards      : (n_env)
88         #TODO 3: Run a step to collect data
89         '''
90         self.mb_states[step, :] = ...
91         self.mb_dones[step, :] = ...
92         self.mb_actions[step, :] = ...
93         self.mb_a_logps[step, :] = ...
94         self.mb_values[step, :] = ...
95         self.states, rewards, self.dones, info = self.env.step(actions)
96         self.mb_rewards[step, :] = ...
97         '''
98
99     last_values = value_net(torch.from_numpy(self.states).float().to(self.device)).cpu().numpy()
100     self.record()
101
```



- **mb_states:** (n_step, n_env, s_dim)
- **mb_actions:** (n_step, n_env, a_dim)
- **mb_dones:** (n_step, n_env)
- **mb_a_logps:** (n_step, n_env)
- **mb_values:** (n_step, n_env)
- **mb_rewards:** (n_step, n_env)

EnvRunner Class (env_runner.py)

Output: $\{s_t, a_t, \log \pi(a_t|s_t), V(s_t), G_t, r(s_t, a_t)\}$



```
102     #2. Compute returns
103     #-----
104     mb_returns = compute_gae(self.mb_rewards, self.mb_values, self.mb_dones, last_values, self.done)
105
106     #mb_states : (n_step*n_env, s_dim)
107     #mb_actions: (n_step*n_env) or (n_env*n_step, a_dim)
108     #mb_a_logps: (n_step*n_env)
109     #mb_values : (n_step*n_env)
110     #mb_returns: (n_step*n_env)
111     return self.mb_states.reshape(self.n_step*self.n_env, self.s_dim), \
112           self.mb_actions.reshape(self.n_step*self.n_env, self.a_dim), \
113           self.mb_a_logps.flatten(), \
114           self.mb_values.flatten(), \
115           mb_returns.flatten()
116
```

PPO Agent

PPO Class

- **PPO** class has the following functions:

#Constructor

- `__init__(self, policy_net, value_net, lr, max_grad_norm, clip_val, sample_n_epoch, sample_mb_size, mb_size, device)`

#Train PPO

- `train(self, mb_states, mb_actions, mb_old_values, mb_advs, mb_returns, mb_old_a_logps)`

#Learning rate decay

- `lr_decay(self, it, n_it)`

PPO Class (agent.py)

- train (TODO 4)

```
33 #Train PPO
34 def train(self, mb_states, mb_actions, mb_old_values, mb_advs, mb_returns, mb_old_a_logps):
35     mb_states = torch.from_numpy(mb_states).to(self.device)
36     mb_actions = torch.from_numpy(mb_actions).to(self.device)
37     mb_old_values = torch.from_numpy(mb_old_values).to(self.device)
38     mb_advs = torch.from_numpy(mb_advs).to(self.device)
39     mb_returns = torch.from_numpy(mb_returns).to(self.device)
40     mb_old_a_logps = torch.from_numpy(mb_old_a_logps).to(self.device)
41
42     for i in range(self.sample_n_epoch):
43         np.random.shuffle(self.rand_idx)
44
45         for j in range(self.sample_n_mb):
46             sample_idx = self.rand_idx[j*self.sample_mb_size : (j+1)*self.sample_mb_size]
47             sample_states = mb_states[sample_idx]
48             sample_actions = mb_actions[sample_idx]
49             sample_old_values = mb_old_values[sample_idx]
50             sample_advs = mb_advs[sample_idx]
51             sample_returns = mb_returns[sample_idx]
52             sample_old_a_logps = mb_old_a_logps[sample_idx]
53
54             sample_a_logps, sample_ents = self.policy_net.evaluate(sample_states, sample_actions)
55             sample_values = self.value_net(sample_states)
56
```

PPO Class (agent.py)

(Hint: Use **sample_a_logps** & **sample_old_a_logps** to compute the probability ratio)

```
57         #PPO loss
58         v_pred_clip = sample_old_values + torch.clamp(sample_values - sample_old_values, -self.clip_value, self.clip_value)
59         v_loss1      = (sample_returns - sample_values).pow(2)
60         v_loss2      = (sample_returns - v_pred_clip).pow(2)
61         v_loss       = torch.max(v_loss1, v_loss2).mean()
62
63         #TODO 4: Policy gradient loss for PPO
64         '''
65         pg_loss = ...
66         '''
67
68         #Train actor
69         self.opt_actor.zero_grad()
70         pg_loss.backward()
71         nn.utils.clip_grad_norm_(self.policy_net.parameters(), self.max_grad_norm)
72         self.opt_actor.step()
73
74         #Train critic
75         self.opt_critic.zero_grad()
76         v_loss.backward()
77         nn.utils.clip_grad_norm_(self.value_net.parameters(), self.max_grad_norm)
78         self.opt_critic.step()
79
80         return pg_loss.item(), v_loss.item()
81
```

Parameters

Parameters (train.py)

```
10 def main():
11     #TODO 5: Adjust these parameters if needed
12     #Parameters that can be modified
13     #-----
14     n_env          = 8
15     n_step         = 128
16     sample_mb_size = 64
17     sample_n_epoch = 4
18     a_std          = 0.5
19     lamb           = 0.95
20     gamma          = 0.99
21     clip_val       = 0.2
22     lr             = 1e-4
23     n_iter         = 30000
24
```


Note:

Too many environments may cause "OSError: The paging file is too small for this operation to complete". If so, please set **n_env** smaller.


- **n_env**: number of environments (actors)
- **n_step**: number of step for runner
- **sample_mb_size**: sample mini-batch size
- **sample_n_epoch**: number of epoch in PPO
- **a_std**: std. dev. of action distribution
- **lamb**: gae factor λ
- **gamma**: discount factor γ
- **clip_val**: clip value ϵ
- **lr**: learning rate
- **n_iter**: number of iteration

Experimental Results

plot.py



play.py



eval.py

```
Total reward = 258.102730, length = 320
Total reward = 206.884708, length = 325
Total reward = 151.192916, length = 307
Total reward = 169.869626, length = 282
Total reward = 150.867508, length = 306
Total reward = 174.895304, length = 298
Total reward = 170.725990, length = 295
Evaluation Score: 159.6902
```

Requirements

- Python3.6+
- Numpy
- Matplotlib
- Opencv
- PyTorch
- Cloudpickle

Execution

- Training: `python train.py`
- Playing: `python play.py`
- Plotting: `python plot.py`
- Evaluating: `python eval.py`

Score & Requirement

- You should complete the code in “model.py”, “env_runner.py”, and “agent.py” and train the model. (TODO 1 ~ 4)
 - Score: TODO 1(15%), TODO 2(15%), TODO 3(15%), TODO 4(15%)
- After training, evaluate your model by executing “plot.py” and “eval.py”.
 - Score: plot(10%), evaluation(30%)
 - Evaluation: $30 * ES / 120$, where ES is your evaluation score (you will get 30 if $ES > 120$)
- Submit the zip of the project folder. It should include:
 - Code (*.py)
 - Training results ("save/")
 - Result folder should include:
 - return record (“return.txt”)
 - Weightings (“model.pt”)

Deadline

- **Deadline: 2022/06/17 (10:00 pm)**
- **Homework Upload**
 - IP: 140.114.79.183 / Port 21
 - Directory: /RNE/HW4
 - Username: RNE_guest
 - Password: RNE2022nthu
 - 請將完整程式碼壓縮成zip並命名為 [學號]_[姓名]_v[版本].zip (如果要更新上傳檔案請設不同的版本號)